

# Chapter 8

# Chapter Precap

- Prior predictive checks and the importance of these for model building.
- How to specify more specific prior probabilities for individual model parameters.
- We introduce heteroscedastic models, that is, models with error terms that vary from observation to observation.
- We present a ‘simple’ model that includes only variation in the error term across two conditions.
- We present a ‘complex’ model that features listener-dependent error terms fit using shrinkage, in addition to the equivalent of listener ‘random effects’ for the error term.
- Finally, we discuss building identifiable models, models supported by the available data, collinearity, linear dependence, and saturated models.

# More About Priors

- In many cases (in linguistics) the prior will have little to no practical effect.
- Use domain knowledge to get priors in the ballpark: 12 lightyears or 12 nanometers?
- Sometimes this is easy (human height), sometimes this is not so easy (the effect of lexical frequency on log reaction times).

# Prior Predictive Checks

- Create fake data sampling only from the prior (ignore the likelihood i.e. the data).

```
# Fit the model yourself
priors = c(brms::set_prior("student_t(3,156, 1000)", class = "Intercept"),
          brms::set_prior("student_t(3,0, 1000)", class = "b"),
          brms::set_prior("student_t(3,0, 1000)", class = "sd"),
          brms::set_prior("lkj_corr_cholesky (1000)", class = "cor"),
          brms::set_prior("student_t(3,0, 1000)", class = "sigma"))

prior_uninformative =
  brms::brm (height ~ A + G + A:G + (A + G + A:G|L) + (1|S),
            ➡ sample_prior="only", data = exp_data, chains = 4, cores = 4,
              warmup = 1000, iter = 5000, thin = 4, prior = priors)
```

# Prior Predictive Checks

Uninformative  
Priors



```
priors = c(brms::set_prior("student_t(3,156, 1000)", class = "Intercept"),  
          brms::set_prior("student_t(3,0, 1000)", class = "b"),  
          brms::set_prior("student_t(3,0, 1000)", class = "sd"),  
          brms::set_prior("lkj_corr_cholesky (1000)", class = "cor"),  
          brms::set_prior("student_t(3,0, 1000)", class = "sigma"))
```

Mildly  
informative  
Priors



```
priors = c(brms::set_prior("student_t(3,156, 12)", class = "Intercept"),  
          brms::set_prior("student_t(3,0, 12)", class = "b"),  
          brms::set_prior("student_t(3,0, 12)", class = "sd"),  
          brms::set_prior("lkj_corr_cholesky (12)", class = "cor"),  
          brms::set_prior("student_t(3,0, 12)", class = "sigma"))
```

Conservative  
Priors



```
priors = c(brms::set_prior("student_t(3,156, 6)", class = "Intercept"),  
          brms::set_prior("student_t(3,0, 6)", class = "b"),  
          brms::set_prior("student_t(3,0, 6)", class = "sd"),  
          brms::set_prior("lkj_corr_cholesky (2)", class = "cor"),  
          brms::set_prior("student_t(3,0, 6)", class = "sigma"))
```

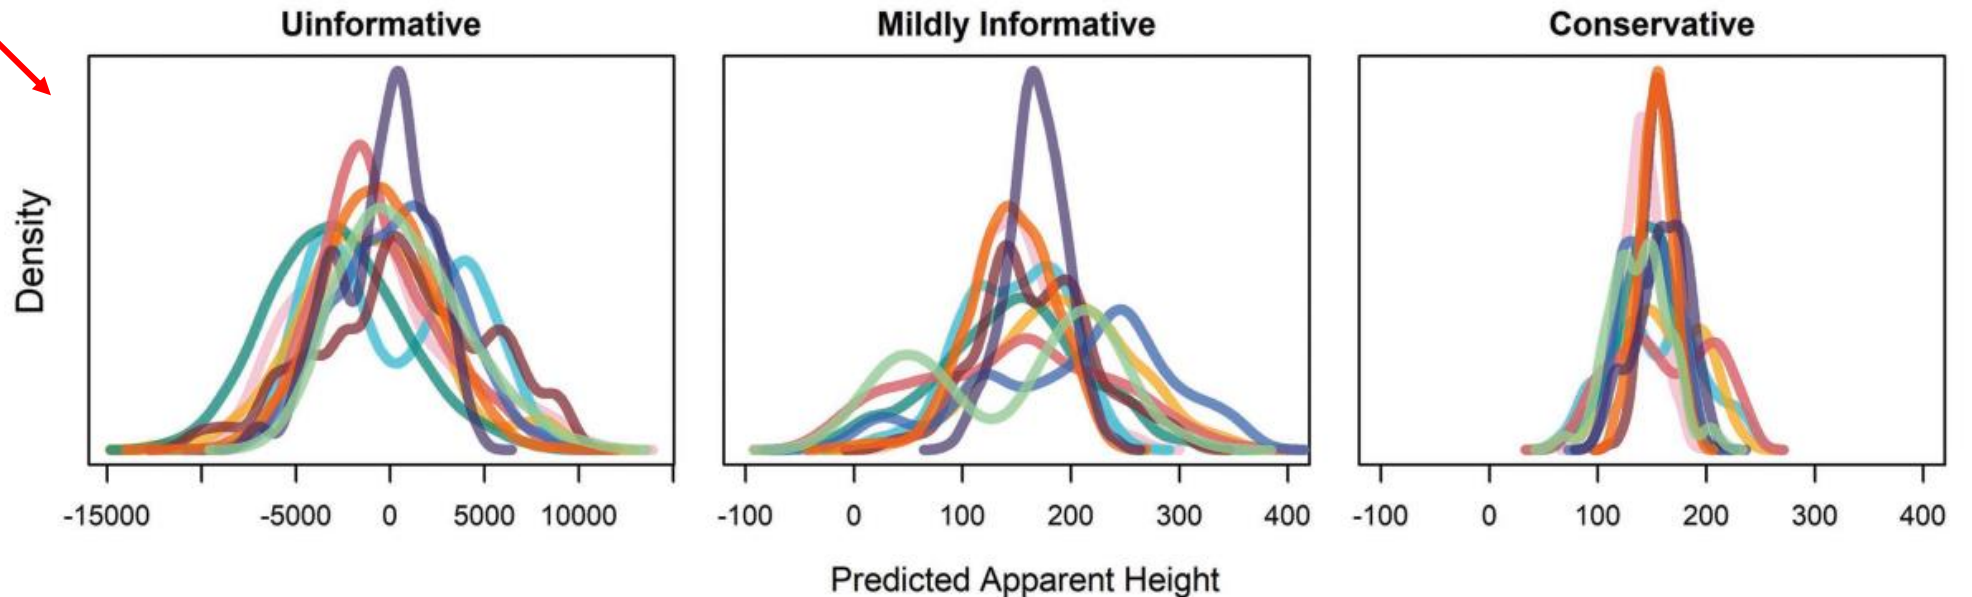
# Prior Predictive Checks

```
pp_uninformative = predict (prior_uninformative, summary = FALSE)
```

```
hist (pp_uninformative[1,])
```

```
bmmb::p_check (pp_uninformative)
```

NOT compared to  
data. Compared  
to domain  
knowledge.

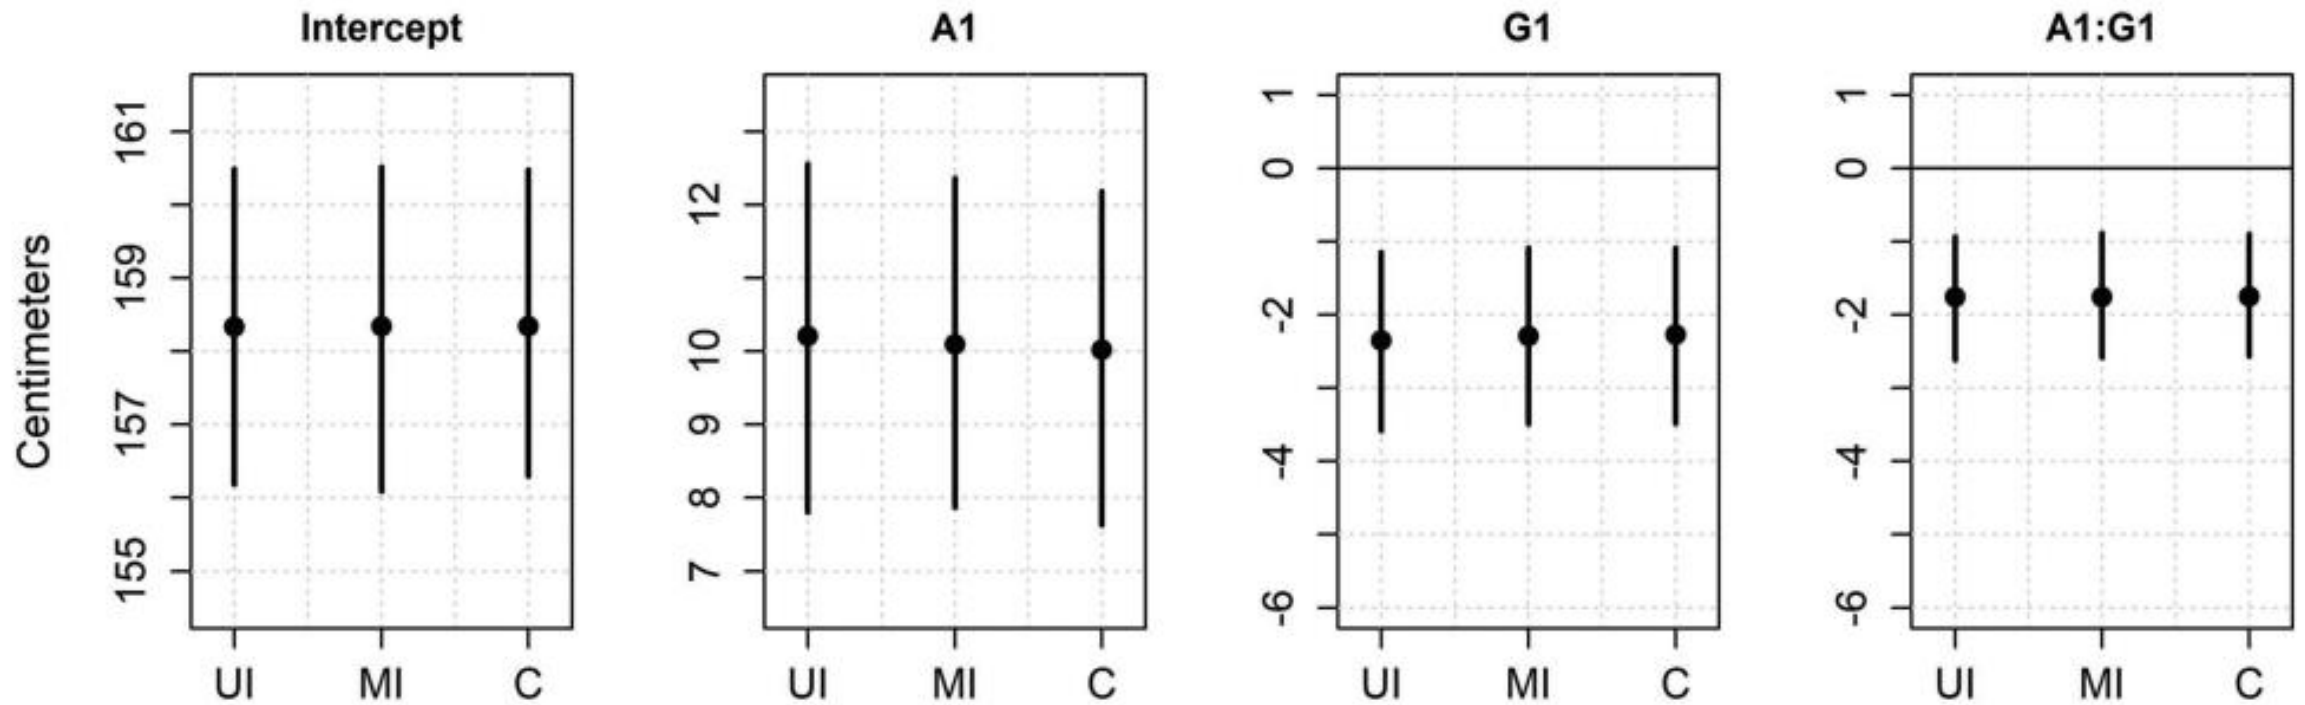


# No Effect on Results!

UI: Uninformative

MI: Mildly informative

C: Conservative



# More Specific Priors

```
# we omit empty columns to let the output fit on the page
bmmmb::prior_summary(model_mildly_informative)[,-c(5:9)]
```

##	prior	class	coef	group	source
##	student_t(3,0, 12)	b			user
##	student_t(3,0, 12)	b	A1		default
##	student_t(3,0, 12)	b	A1:G1		default
##	student_t(3,0, 12)	b	G1		default
##	student_t(3,156, 12)	Intercept			user
##	lkj_corr_cholesky (2)	L			user
##	lkj_corr_cholesky (2)	L		L	default
##	student_t(3,0, 12)	sd			user
##	student_t(3,0, 12)	sd		L	default
##	student_t(3,0, 12)	sd	A1	L	default
##	student_t(3,0, 12)	sd	A1:G1	L	default
##	student_t(3,0, 12)	sd	G1	L	default
##	student_t(3,0, 12)	sd	Intercept	L	default
##	student_t(3,0, 12)	sd		S	default
##	student_t(3,0, 12)	sd	Intercept	S	default
##	student_t(3,0, 12)	sigma			user



# More Specific Priors

```
priors =  
  c(set_prior("student_t(3,156, 6)", class = "Intercept"),  
    set_prior("student_t(3,0, 6)", class = "b"),  
    set_prior("student_t(3,0, 10)", class = "b", coef = "A1"),  
    set_prior("student_t(3,0, 3)", class = "b", coef = "G1"),  
    set_prior("student_t(3,0, 10)", class = "sd"),  
    set_prior("student_t(3,0, 5)", class = "sd", coef = "A1", group="L"),  
    set_prior("student_t(3,0, 1.5)", class = "sd", coef = "G1", group="L"),  
    set_prior("lkj_corr_cholesky (2)", class = "cor"),  
    set_prior("student_t(3,0, 6)", class = "sigma"))
```

```
bmmb::prior_summary(prior_informative)[,-c(5:8)]  
##           prior      class      coef group  
## student_t(3,0, 6)      b  
## student_t(3,0, 10)      b      A1  
## student_t(3,0, 10)      b    A1:G1  
## student_t(3,0, 3)      b      G1  
## student_t(3,156, 6) Intercept  
## lkj_corr_cholesky (2)      L  
## lkj_corr_cholesky (2)      L      L  
## student_t(3,0, 10)      sd  
## student_t(3,0, 10)      sd      L  
## student_t(3,0, 5)      sd      A1      L  
## student_t(3,0, 5)      sd    A1:G1      L  
## student_t(3,0, 1.5)      sd      G1      L  
## student_t(3,0, 1.5)      sd Intercept  L  
## student_t(3,0, 1.5)      sd      S  
## student_t(3,0, 1.5)      sd Intercept  S  
## student_t(3,0, 6)      sigma
```

# Data and Research Questions

```
library (brms)
library (bmmb)
data (exp_data)
options (contrasts = c('contr.sum', 'contr.sum'))
```

- **L**: A number from 1 to 15 indicating which *listener* responded to the trial.
- **height**: A number representing the *height* (in centimeters) reported for the speaker on each trial.
- **S**: A number from 1 to 139 indicating which *speaker* produced the trial stimulus.
- **G**: The *apparent gender* of the speaker indicated by the listener, **f** (female) or **m** (male).
- **A**: The *apparent age* of the speaker indicated by the listener, **a** (adult) or **c** (child).

(Q1) Does our error standard deviation vary as a function of apparent speaker age?

(Q2) Does our error standard deviation vary as a function of the listener?

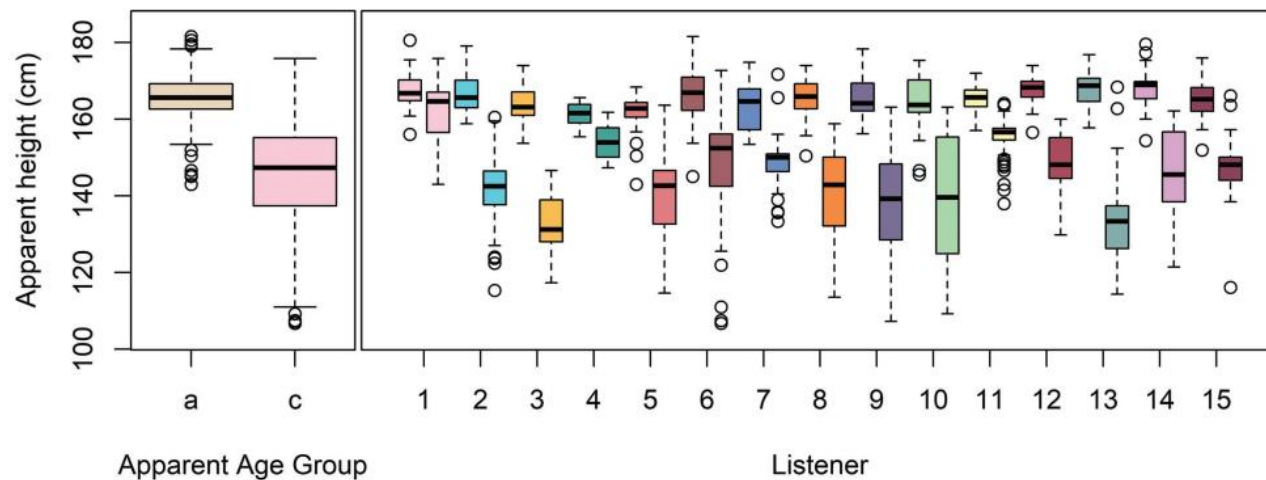
# Homoskedasticity

- Most regression models only predict variation in means.
- A single error variance is assumed for all data (i.e.,  $\sigma$  doesn't get a subscript).

$$\text{height}_{[i]} \sim \mathcal{N}(\mu_{[i]}, \sigma)$$

↓

$$\mu_{[i]} = x_1 + x_2 + \dots + x_p$$



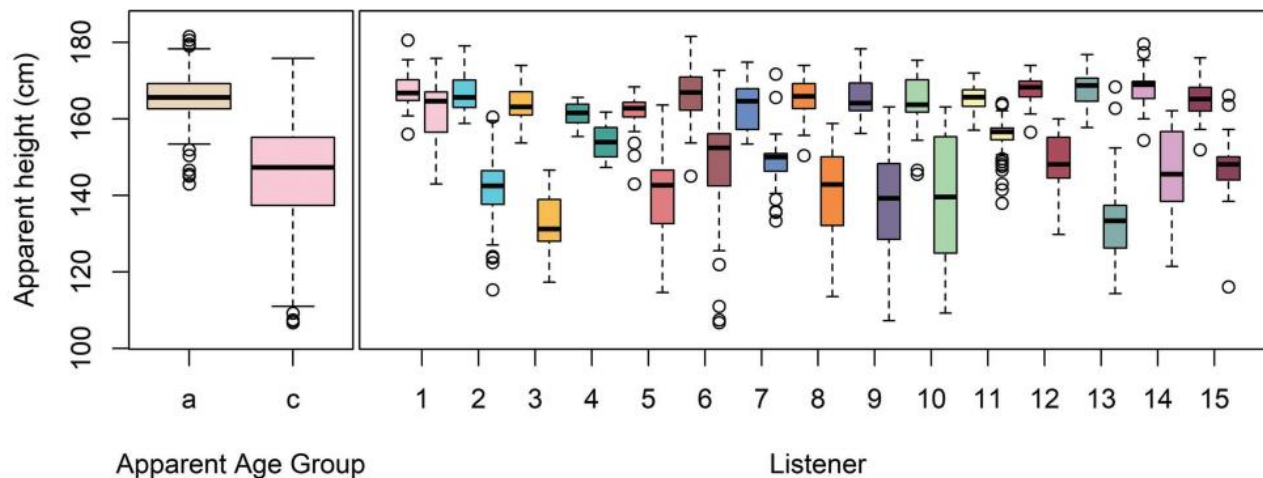
# Heteroskedasticity

- With Bayesian models it is trivial to model variation in error variances as well.
- Different error variances can exist for different data (i.e.,  $\sigma$  does get a subscript).

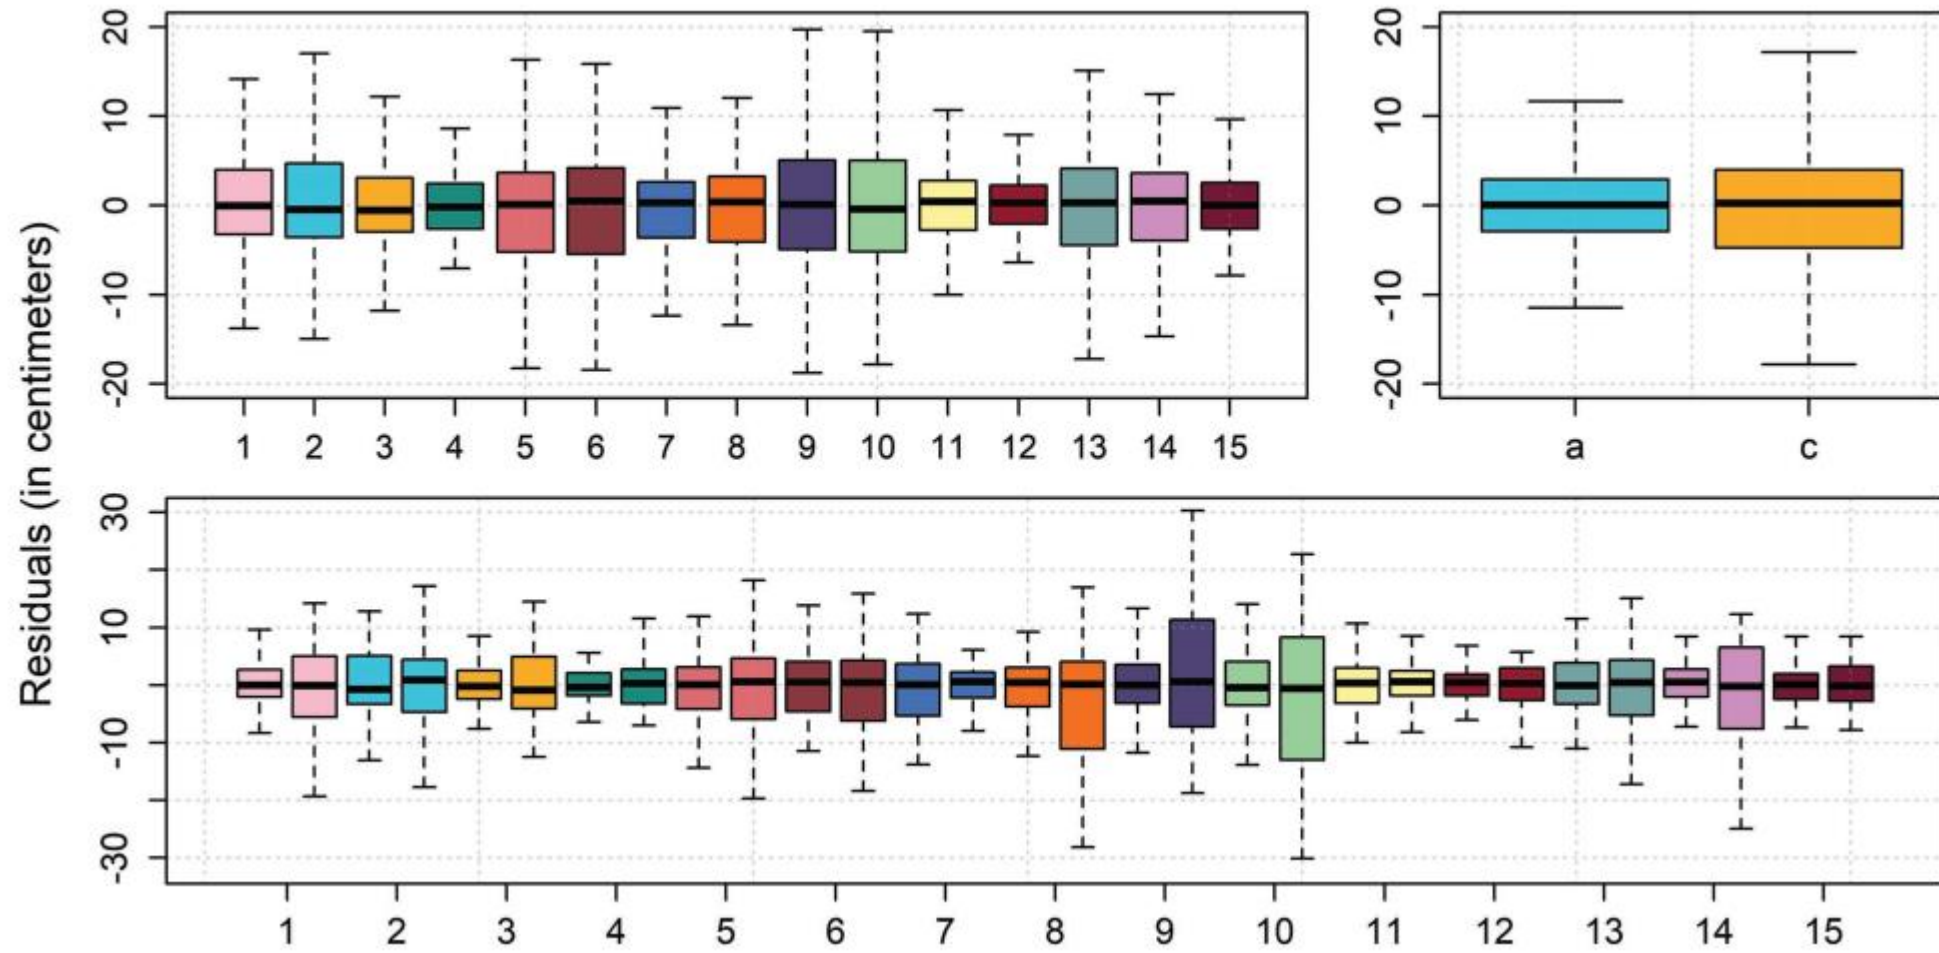
$$\text{height}_{[i]} \sim \mathcal{N}(\mu_{[i]}, \sigma_{[i]})$$

$$\mu_{[i]} = x_1 + x_2 + \dots + x_p$$

$$\sigma_{[i]} = x_{\sigma 1} + x_{\sigma 2} + \dots + x_{\sigma p}$$



# Heteroskedasticity



# Description of Our Model

Chapter 6

→ `height ~ A + G + A:G + (A + G + A:G | L) + (1 | S)`

New component,  
similar to chapter 5.

→ `sigma ~ A`

Together in one model

→  
`model_formula = brms::bf(height ~ A*G + (A*G | L) + (1 | S),  
sigma ~ A + (A | L))`

# Description of Our Model

$$\begin{aligned} \text{height}_{[i]} &\sim t(v, \mu_{[i]}, \sigma_{[i]}) \\ \mu_{[i]} &= \text{Intercept} + A + G + A:G + \\ &L_{[L[i]]} + A:L_{[L[i]]} + G:L_{[L[i]]} + A:G:L_{[L[i]]} + S_{[S[i]]} \end{aligned}$$

$$\log(\sigma_{[i]}) = \text{Intercept}_{\sigma} + A_{\sigma}$$

Priors:

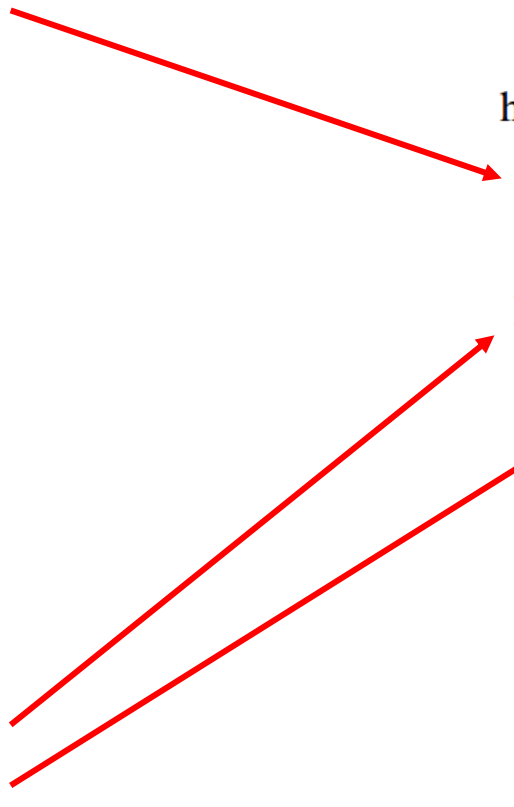
$$S_{[\cdot]} \sim t(3, 0, \sigma_S)$$

$$\begin{bmatrix} L_{[\cdot]} \\ A:L_{[\cdot]} \\ G:L_{[\cdot]} \\ A:G:L_{[\cdot]} \end{bmatrix} \sim \text{MVNormal} \left( \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \Sigma \right)$$




$$\begin{aligned} \text{Intercept} &\sim t(3, 156, 12) \\ A, G, A:G &\sim t(3, 0, 12) \\ \sigma_L, \sigma_{A:L}, \sigma_{G:L}, \sigma_{A:G:L}, \sigma_S &\sim t(3, 0, 12) \\ v &\sim \text{gamma}(2, 0.1) \\ R &\sim \text{LKJCorr}(2) \\ \text{Intercept}_{\sigma} &\sim N(0, 1.5) \\ A_{\sigma} &\sim N(0, 1.5) \end{aligned}$$

$$\begin{aligned} \text{height}_{[i]} &\sim t(v, \mu_{[i]}, \sigma_{[i]}) \\ \sigma_{[i]} &= \text{Intercept}_{\sigma} + A_{\sigma} \end{aligned}$$

$$\begin{aligned} \text{Intercept}_{\sigma} &\sim t(3, 0, 1.5) \\ A_{\sigma} &\sim t(3, 0, 1.5) \end{aligned}$$




# Fitting the Model

```
model_formula = brms::bf(height ~ A*G + (A*G|L) + (1|S),  
                          sigma ~ A)   
  
priors =  
  c(set_prior("student_t(3, 156, 12)", class = "Intercept"),  
    set_prior("student_t(3, 0, 12)", class = "b"),  
    set_prior("student_t(3, 0, 12)", class = "sd"),  
    set_prior("gamma(2, 0.1)", class = "nu"),  
    set_prior("normal(0, 1.5)", class = "Intercept", dpar = "sigma"),   
    set_prior("normal(0, 1.5)", class = "b", dpar = "sigma"),  
    set_prior("lkj_corr_cholesky(2)", class = "cor"))   
  
prior_A_sigma =  
  brms::brm(model_formula, data = exp_data, chains = 4, cores = 4,  
            warmup = 1000, iter = 3500, thin = 2, family="student",  
            prior = priors, sample_prior = "only")
```



# Inspecting the Fixed Effects

```
fixef (model_A_sigma)
##              Estimate Est.Error      Q2.5      Q97.5
## Intercept      158.383    1.10922  156.193  160.6196
## sigma_Intercept    1.724    0.03265    1.659    1.7881
## A1              11.271    1.19021    8.915   13.6901
## G1              -3.067    0.59091   -4.271   -1.9084
## A1:G1            -1.528    0.43464   -2.403   -0.6654
## sigma_A1         -0.247    0.02199   -0.290   -0.2036
```



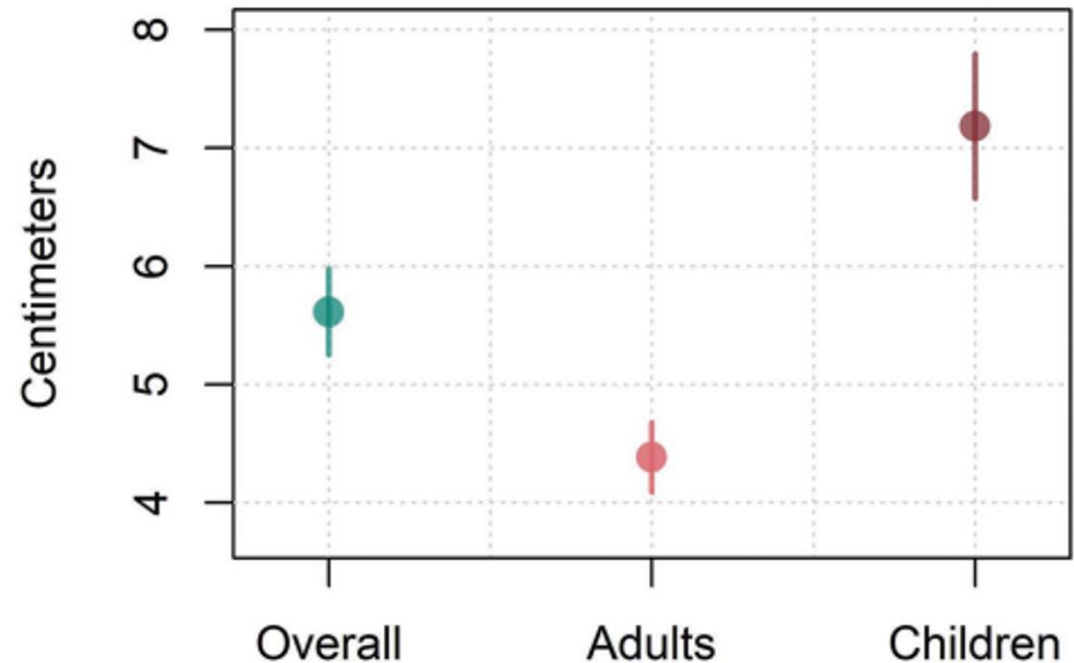
In log units. Must exponentiate values to get original units, i.e.,  $\sigma = \exp(\text{sigma\_Intercept})$

# Recovering predicted Sigmas

You MUST combine parameters before exponentiating.  
Otherwise, the results will be totally wrong.

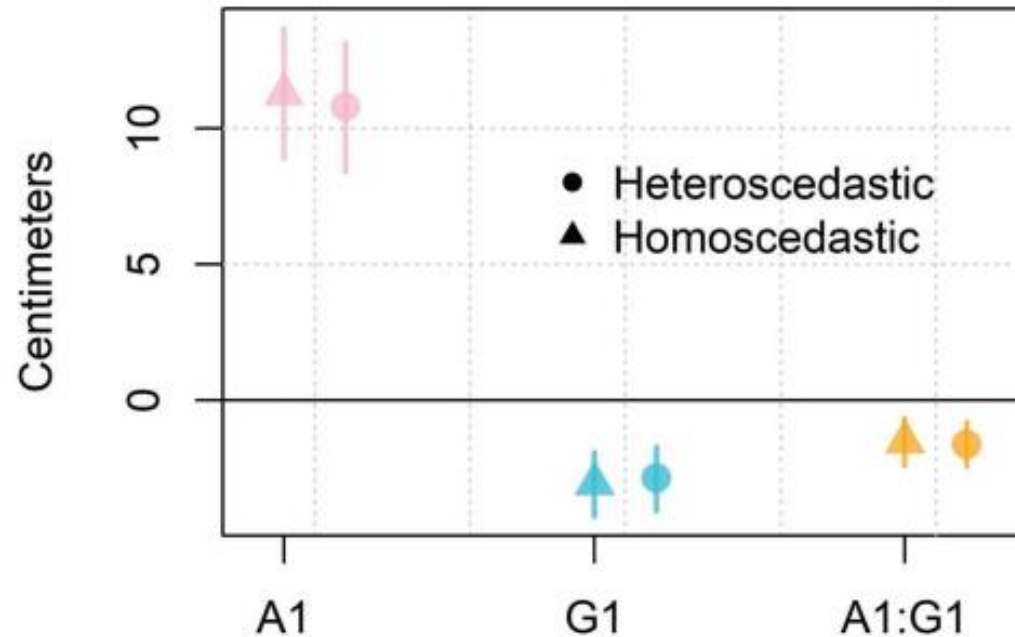
```
sigmas = short_hypothesis(  
  model_A_sigma,  
  c("exp(sigma_Intercept) = 0",           # overall sigma  
    "exp(sigma_Intercept + sigma_A1) = 0", # adult sigma  
    "exp(sigma_Intercept - sigma_A1) = 0")) # child sigma
```

```
sigmas[, -5]  
##      Estimate Est.Error  Q2.5  Q97.5  
## H1      5.609    0.1830  5.253  5.978  
## H2      4.382    0.1526  4.085  4.680  
## H3      7.184    0.3115  6.573  7.793
```



# Model Comparison

```
model_interaction = add_criterion(model_interaction, "loo")
model_A_sigma = add_criterion(model_A_sigma, "loo")
loo_compare (model_interaction, model_A_sigma)
##                               elpd_diff se_diff
## model_A_sigma                 0.0         0.0
## model_interaction            -53.7        12.8
```



# A 'Complex' Model

- Our last model had only a 'fixed effect' for apparent age.

```
sigma ~ A
```

- If we were predicting the mean, we would want to include random effects for listener.

- We can include these in our prediction of sigma as well.

```
sigma ~ A + (A|L)
```

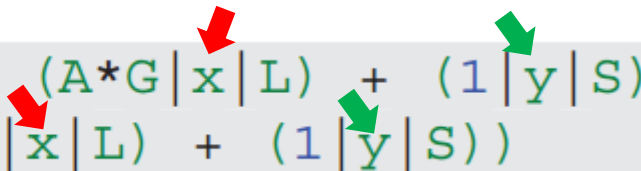
# Our Model Formula

- When we have 'the same' random effects across formulas, we need to let our model know this.

```
model_formula = brms::bf(height ~ A*G + (A*G|L) + (1|S),  
                           sigma ~ A + (A|L))
```

- We do this by putting the same indicator between pipes (|inhere|) before a random effect.

```
model_formula = brms::bf(height ~ A*G + (A*G|x|L) + (1|y|S),  
                           sigma ~ A + (A|x|L) + (1|y|S))
```



- Our model will estimate the correlation of random effects across predicted variables!

# Model Description

$$\begin{aligned} \text{height}_{[i]} &\sim \text{N}(\mu_{[i]}, \sigma_{[i]}) \\ \mu_{[i]} &= \text{Intercept} + A + G + A : G + \\ &L_{[L[i]]} + A : L_{[L[i]]} + G : L_{[L[i]]} + A : G : L_{[L[i]]} + S_{[S[i]]} \\ \log(\sigma_{[i]}) &= \text{Intercept}_{\sigma} + A_{\sigma} + A_{\sigma} : L_{\sigma[L[i]]} + L_{\sigma[L[i]]} \end{aligned}$$

Priors :

$$\begin{bmatrix} L_{[\cdot]} \\ A : L_{[\cdot]} \\ G : L_{[\cdot]} \\ A : G : L_{[\cdot]} \\ L_{\sigma[\cdot]} \\ A_{\sigma} : L_{\sigma[\cdot]} \end{bmatrix} \sim \text{MVNormal} \left( \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \Sigma \right)$$

$$S_{[\cdot]} \sim \text{N}(0, \sigma_S)$$

$$\text{Intercept} \sim \text{t}(3, 156, 12)$$

$$A \sim \text{t}(3, 0, 12)$$

$$\sigma, \sigma_L, \sigma_{A:L}, \sigma_S \sim \text{t}(3, 0, 12)$$

$$\text{Intercept}_{\sigma} \sim \text{N}(0, 1.5)$$

$$A_{\sigma} \sim \text{N}(0, 1.5)$$

$$\sigma_{L_{\sigma}}, \sigma_{A_{\sigma}:L_{\sigma}} \sim \text{N}(0, 1.5)$$

$$R \sim \text{LKJCorr}(2)$$

# Fitting the Model

```
# Fit the model yourself
model_formula = brms::bf(height ~ A*G + (A*G|x|L) + (1|S),
                          sigma ~ A + (A|x|L))

priors =
  c(set_prior("student_t(3, 156, 12)", class = "Intercept"),
    set_prior("student_t(3, 0, 12)", class = "b"),
    set_prior("student_t(3, 0, 12)", class = "sd"),
    set_prior("gamma(2, 0.1)", class = "nu"),
    set_prior("normal(0, 1.5)", class = "Intercept", dpar = "sigma"),
    set_prior("normal(0, 1.5)", class = "b", dpar = "sigma"),
    set_prior("normal(0, 1.5)", class = "sd", dpar = "sigma"),
    set_prior("lkj_corr_cholesky (2)", class = "cor"))

model_A_L_sigma =
  brms::brm (model_formula, data = exp_data, chains = 4, cores = 4,
            warmup = 1000, iter = 3500, thin = 2, family="student",
            prior = priors)
```

# Inspecting the Results

```
bmbb::short_summary (model_A_L_sigma)
## ...
## sd(sigma_Intercept)      0.36      0.08      0.24      0.54
## sd(sigma_A1)             0.17      0.04      0.10      0.27
## ...
## cor(A1,sigma_A1)         -0.29      0.22     -0.67      0.16
## cor(G1,sigma_A1)         0.22      0.24     -0.27      0.65
## cor(A1:G1,sigma_A1)      0.12      0.24     -0.35      0.58
## cor(sigma_Intercept,sigma_A1) -0.44      0.21     -0.79      0.05
## ...
##
## Population-Level Effects:
##               Estimate Est.Error l-95% CI u-95% CI
## Intercept      158.29      1.10    156.15    160.51
## sigma_Intercept  1.74      0.10     1.55     1.93
## A1             11.28      1.18     8.98     13.64
## G1             -2.92      0.57    -4.06    -1.80
## A1:G1           -1.63      0.42    -2.47    -0.77
## sigma_A1       -0.23      0.05    -0.33    -0.14
##
## Family Specific Parameters:
##      Estimate Est.Error l-95% CI u-95% CI
## nu      8.07      1.54     5.71    11.72
```

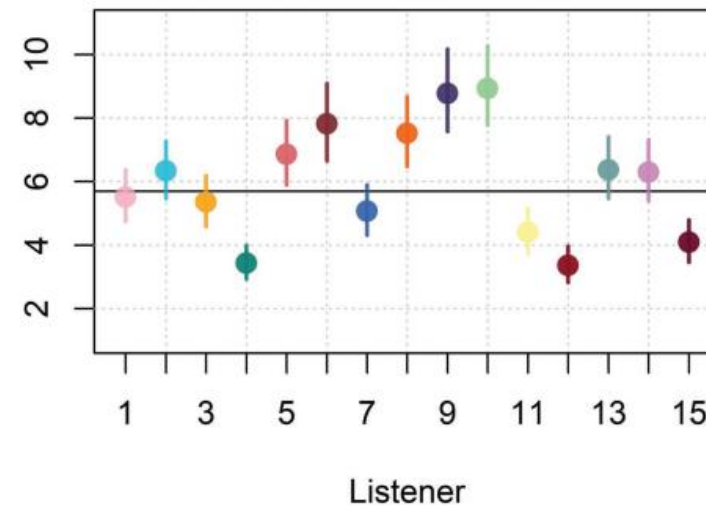
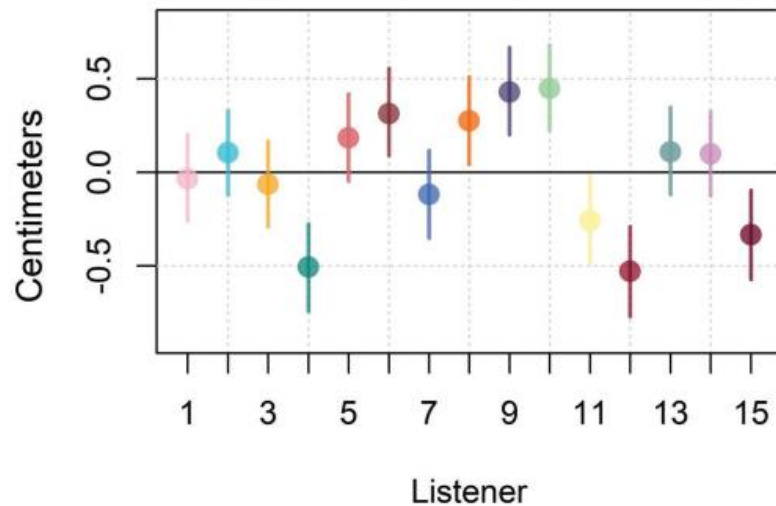


# Listener-dependent Error Terms

```
# listener random effects for sigma
log_sigmas_centered = short_hypothesis(model_A_L_sigma, "sigma_Intercept=0",
                                       scope = "ranef", group="L")

# the sum of the sigma intercept and the listener sigma random effect
log_sigmas = short_hypothesis(model_A_L_sigma, "sigma_Intercept=0",
                              scope = "coef", group="L")

# the exponent of the sum of the sigma intercept
# and the listener sigma random effect
sigmas = short_hypothesis(model_A_L_sigma, "exp(sigma_Intercept)=0",
                          scope = "coef", group="L")
```



# Model Comparison

- Our new model is much better.

```
model_A_L_sigma = add_criterion(model_A_L_sigma, "loo")
```

```
loo_compare (model_interaction, model_A_sigma, model_A_L_sigma)
##               elpd_diff se_diff
## model_A_L_sigma      0.0      0.0
## model_A_sigma     -122.1     15.6
## model_interaction -175.8     20.8
```

# Answering our Research Questions

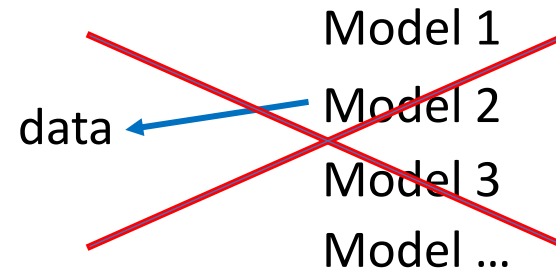
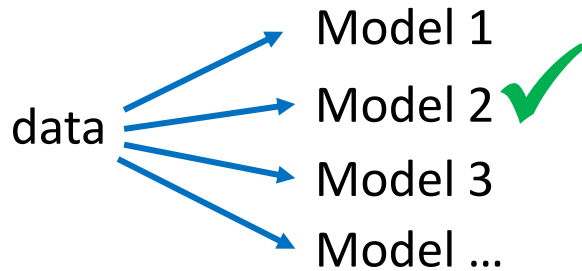
(Q1) Does our error standard deviation vary as a function of apparent speaker age?

(Q2) Does our error standard deviation vary as a function of the listener?

- Yes, and yes.
- Some tougher questions are: Which is the 'better/rea' model? Which model should we report?

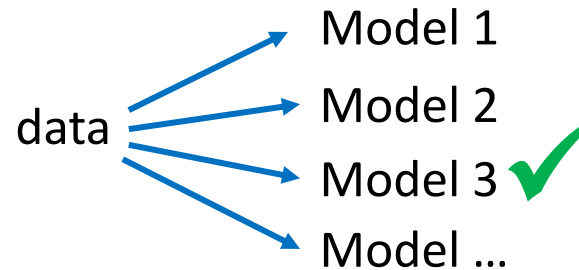
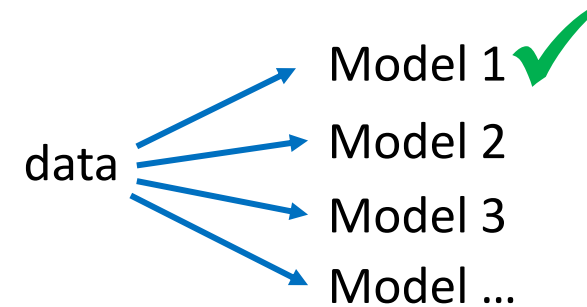
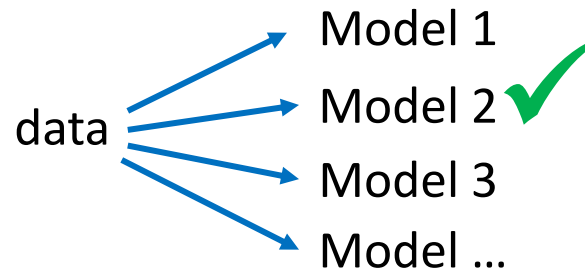
# Underdetermination

- Given some observations, there are always multiple different interpretations/models.



# Researcher Degrees of Freedom

- The decisions a researcher makes during design or analysis that affect the model used and conclusions reached.

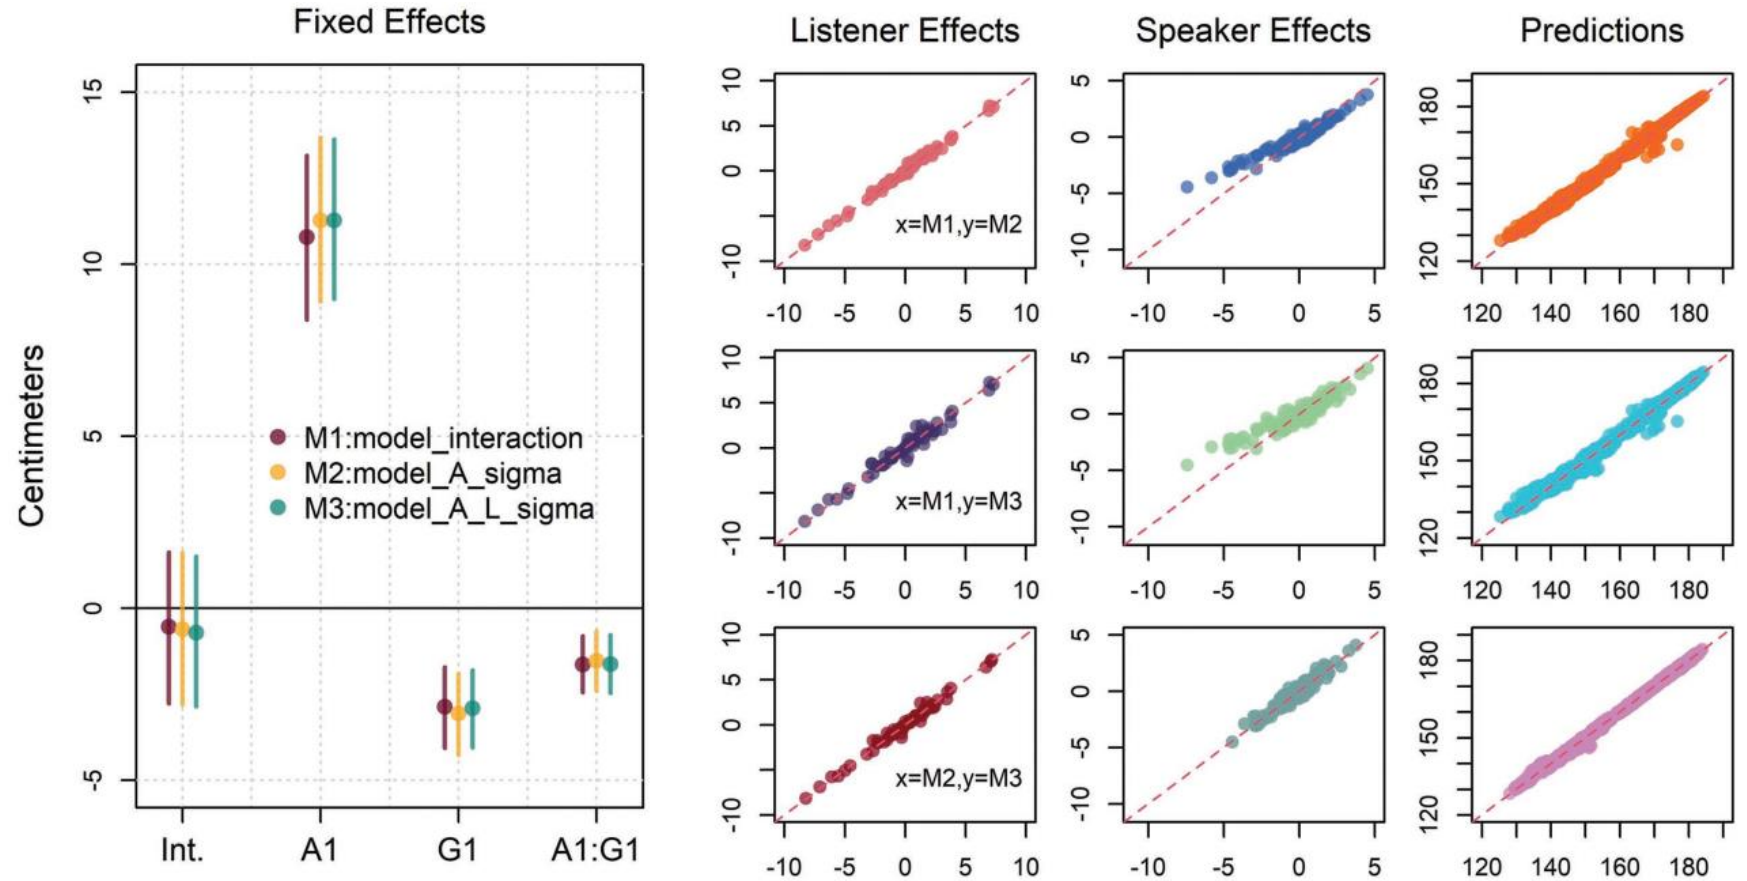


# Which Model To Report?

- Our models represent different *information*.
- When considering which to report, we can consider two things:
  1. Which parameters/information are we interested in?
  2. How much do these differ across models?
    - 2a. If they differ, why?

# Which Model To Report?

- Fixed effects are about the same.
- Speaker effects differ in heteroskedastic models.
- Do you care about the variation in the errors?



# Should You fit the Model?

- As we progress, you *can* fit more and more complicated models.
- But *should* you?
- Two things to worry about:
  - Identifiability: the ability to estimate unique, independent values for all your parameters.
  - Support: Having enough data to realistically estimate all your parameters.



# Collinearity

- A set of predictor vectors is linearly independent when there is no vector of non-zero numbers that can be used to combine our predictors such that they always equal zero.

$$0 = x_1 \cdot a_1 + x_2 \cdot a_2 + \dots + x_n \cdot a_n$$



If this is possible the x predictors are not linearly dependent.

# Collinearity

- You cannot fit a model using height in meters and height in centimeters, and estimate both effects independently.

```
exp_data$vtl_m = exp_data$vtl / 100
```

```
exp_data$vtl + exp_data$vtl_m * -100
```

- These predictors are linearly dependent!

$$x_1 \cdot a_1 + x_2 \cdot a_2 = 0$$

$$\text{vtl} \cdot 1 + \text{vtl}_m \cdot -100 = 0$$

# Collinearity

```
model_bad_1 =  
  brms::brm (height ~ vtl_m + vtl, data = exp_data, chains = 4, cores = 4,  
    warmup = 1000, iter = 3500, thin = 2,  
    prior = c(brms::set_prior("normal(176, 50)", class = "Intercept"),  
      brms::set_prior("normal(0, 15)", class = "sigma")))
```

## Population-Level Effects:

##		Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
##	Intercept	45.69	2.37	41.17	50.47	1.00	4031	4425
##	vtl_m	1383.91	246800.93	-531293.80	343951.54	2.12	5	13
##	vtl	-5.28	2468.01	-3430.88	5321.56	2.12	5	13

# Semi-Collinearity

- Linear dependence is binary.
- Correlation is a continuous measure of linear dependence (basically).

```
cor (exp_data$vtl_m, exp_data$vtl)
## [1] 1
```

- What if we make our predictors almost linearly dependent?

```
set.seed(1)
exp_data$vtl_m_noise = exp_data$vtl_m +
  rnorm (length(exp_data$vtl_m), 0, sd(exp_data$vtl_m)/10)

cor (exp_data$vtl, exp_data$vtl_m_noise)
## [1] 0.9946
```

# Semi-Collinearity: Not that Bad

```
model_bad_2 =  
  brms::brm (height ~ vtl_m_noise + vtl, data = exp_data, chains = 4,  
            cores = 4, warmup = 1000, iter = 3500, thin = 2, prior = priors)
```

## Population-Level Effects:

##	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
## Intercept	45.75	2.36	41.16	50.35	1.00	4523	4087
## vtl_m_noise	-201.86	175.11	-544.11	139.03	1.00	3504	3697
## vtl	10.57	1.76	7.14	14.01	1.00	3439	3866

# Semi-Collinearity: Could be Better

```
model_good =  
  brms::brm (height ~ vtl, data = exp_data, chains = 4, cores = 4,  
    warmup = 1000, iter = 3500, thin = 2,  
    prior = c(brms::set_prior("normal(176, 50)", class = "Intercept"),  
      brms::set_prior("normal(0, 15)", class = "sigma")))
```

## Population-Level Effects:

##		Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
##	Intercept	45.72	2.38	41.04	50.40	1.00	5052	4679
##	vtl	8.55	0.18	8.21	8.90	1.00	5072	4831

```
model_bad_2 =
```

```
  brms::brm (height ~ vtl_m_noise + vtl, data = exp_data, chains = 4,  
    cores = 4, warmup = 1000, iter = 3500, thin = 2, prior = priors)
```

## Population-Level Effects:

##		Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
##	Intercept	45.75	2.36	41.16	50.35	1.00	4523	4087
##	vtl_m_noise	-201.86	175.11	-544.11	139.03	1.00	3504	3697
##	vtl	10.57	1.76	7.14	14.01	1.00	3439	3866

# Predictable Values for Predictors

- You can't fit models where the value of one categorical predictors can be guessed based on the values of other predictors.
- This is again a problem of linear dependence.
- This is why we can't estimate both levels of a two-group factor.

```
x = cbind (intercept=rep(1,4), x1=rep(c(1,0),2),  
x2=rep(c(0,1),2))  
x  
##           intercept x1 x2  
## [1,]             1  1  0  
## [2,]             1  0  1  
## [3,]             1  1  0  
## [4,]             1  0  1
```

```
x[,1] + x[,2]*(-1) + x[,3]*(-1)  
## [1] 0 0 0 0
```

# Predictable Values for Predictors

- And also why we can't estimate all levels of a four level factor.

```
x = cbind (intercept=rep(1,4), C1=c(1,0,0,0), C2=c(0,1,0,0),  
           C3=c(0,0,1,0), C4=c(0,0,0,1))
```

```
X
```

```
##      intercept C1 C2 C3 C4  
## [1,]         1  1  0  0  0  
## [2,]         1  0  1  0  0  
## [3,]         1  0  0  1  0  
## [4,]         1  0  0  0  1
```

```
x[,1] + x[,2]*(-1) + x[,3]*(-1) + x[,4]*(-1) + x[,5]*(-1)  
## [1] 0 0 0 0
```



# Predictable Values for Predictors

- And also why we can't include group, age, *and* gender.

```
x = cbind (intercept=rep(1,4), C1=c(1,0,0,0), C2=c(0,1,0,0),  
          C3=c(0,0,1,0), A1=c(0,0,1,1), G1=c(0,1,0,1), A1G1=c(1,0,0,1))
```

```
X  
##      intercept C1 C2 C3 A1 G1 A1G1  
## [1,]         1  1  0  0  0  0    1  
## [2,]         1  0  1  0  0  1    0  
## [3,]         1  0  0  1  1  0    0  
## [4,]         1  0  0  0  1  1    1
```

```
x[,1]*1 + x[,2]*(-1) + x[,3]*(-1) + x[,5]*(-1)  
## [1] 0 0 0 0
```

```
x[,1]*1 + x[,3]*(-1) + x[,4]*(-1) + x[,7]*(-1)  
## [1] 0 0 0 0
```

# Predictable Values for Predictors

```
model_bad_3 =  
  brms::brm (height ~ C + A*G, data = exp_data, chains = 4, cores = 4,  
    warmup = 1000, iter = 3500, thin = 2,  
    prior = c(brms::set_prior("normal(176, 15)", class = "Intercept"),  
      brms::set_prior("normal(0, 15)", class = "sigma")))
```

## Population-Level Effects:

##	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk ESS	Tail ESS
## Intercept	157.98	0.20	157.57	158.35	1.06	61	78
## C1	1066.31	2604.72	-4431.79	5158.68	2.07	5	18
## C2	496.91	1751.32	-2543.19	3708.83	1.87	6	13
## C3	74.39	2086.76	-4113.31	2801.74	1.99	5	21
## A1	793.68	1363.24	-2383.78	3019.96	1.90	6	12
## G1	567.63	2171.69	-4107.90	3364.23	2.04	5	18
## A1:G1	283.89	1112.60	-2295.30	2141.73	2.26	5	18

This is not ideal.

# Saturated Models

- Saturated models have one parameter for every observation.
- Without shrinkage this means that there is no random variation in the model, i.e., the error cannot be estimated.

```
exp_data$S = factor (exp_data$S)
exp_data$L = factor (exp_data$L)

model_bad_4 =
  brms::brm (height ~ S*L, data = exp_data, chains = 4, cores = 4,
    warmup = 1000, iter = 3500, thin = 2,
    prior = c(brms::set_prior("normal(176, 15)", class = "Intercept"),
      brms::set_prior("normal(0, 15)", class = "sigma")))
```



This model crashes my session of R.

# Nearly-Saturated Models

- A model can be 'nearly'-saturated.
- What if you have 2 observations per speaker per listener. You can fit this model (probably), but should you?
- Follow up: What is your  $n$ ?
- Think of your  $n$  for each individual parameter rather than overall.

# Exercises: Week 7

Use the data in 'exp\_ex' to do one of the following. You may also use your own data to answer a related question. In any case, describe the model, present and explain the results. More requirements:

- You must include a model at least this complex:  $y \sim A*B + (A*B|L)$ , meaning two factors and an interaction.
- Fit a new model, not like in the book and not like for week 6.
- Include at least two non-book figures.
- Recreate the predicted group means and report them.
- Report and interpret at least one simple effect.